

**ISSRE 2019**  
Tutorial Session

# Building Applications for Trustworthy Data Analysis in the Cloud

Andrey Brito  
André Martin  
Lilia Sampaio  
Fábio Silva

# Security-aware data processing

Part 1



Why secure  
data processing?



1  
Users want data  
to be processed  
in the cloud

In 2019, companies  
executed

**79%**

of their workload in  
the cloud

(RightScale 2019 - State of the  
Cloud Report)

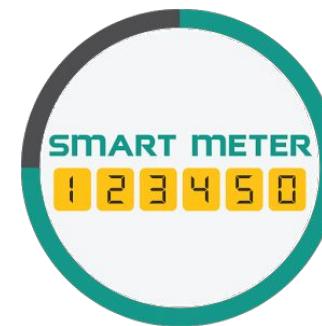
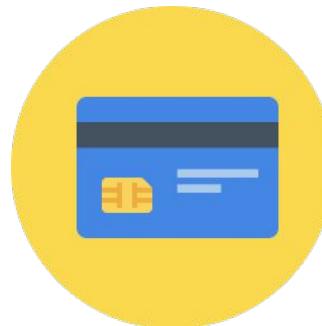
Up to 2021,

**94%**

of this workload  
will be processed in  
the cloud

(Cisco Global Cloud Index:  
Forecast and Methodology,  
2016-2021 White Paper)

# Sensitive data requires increasing the level of security measures when processing and storing such data



Personal  
information

Health

Financial

Energy  
consumption

Company  
strategy related



1

Users want data  
to be processed  
in the cloud



2

Security is in  
the TOP 5 cloud  
challenges,  
being cited for  
over 81% of  
participants

During the first

6 MONTHS OF 2018

the equivalent of

291 RECORDS



was stolen or exposed

EVERY SECOND!

(Source: [Article “2018: The year of the data breach tsunami”](#) - MalwarebytesLABS, 2018)



1 Users want data to be processed in the cloud



2 Security is in the TOP 5 cloud challenges, being cited for over 81% of participants

Secure data processing is then very important!



# How to securely process sensitive data?

# Intel SGX

Software Guard eXtensions

- Trusted execution environments
  - Hardware technology
  - Guarantees of data integrity and confidentiality
  - Use of isolated and protected memory areas called *enclaves*
  - Attestation process
-

# SCONE

Secure CONtainer Environment

- Uses SGX to protect container processes
- Transparent to already existing Docker environments
- There are no changes to the application code being deployed
- Prepares the code to be SGX-compatible



## Top Cloud Initiatives in 2019

**64%**

Optimize existing  
use of cloud

**58%**

Move more  
workloads to cloud

**39%**

Expand use of  
containers



**1**

Users want data  
to be processed  
in the cloud



**2**

Security is in  
the TOP 5 cloud  
challenges,  
being cited for  
over 81% of  
participants

Secure data  
processing is then  
very important!



**3**

Resources should be  
managed in order to  
attend users needs



1

Users want data  
to be processed  
in the cloud



2

Security is in  
the TOP 5 cloud  
challenges,  
being cited for  
over 81% of  
participants

Secure data  
processing is then  
very important!

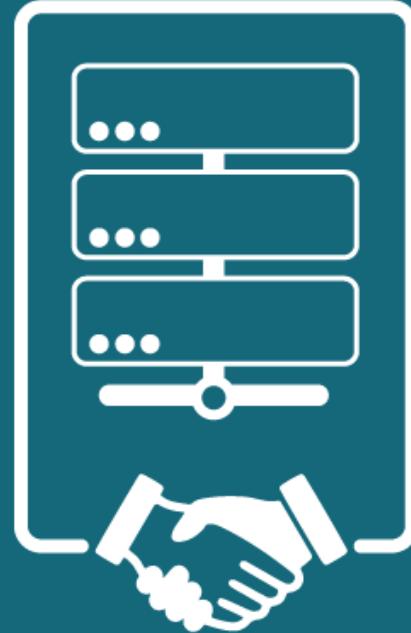


3

Resources should be  
managed in order to  
attend users needs

4

Quality of Service!



# QoS and Reliability

Quality of Service as a reliability measure

- QoS management can be defined as "*the allocation of resources to an application in order to guarantee a service level along dimensions such as performance, availability and reliability*"

(Ardagna et al. (2014) - *Quality-of-service in cloud computing: modeling techniques and their applications*)

— — —



Cloud support

Data processing

Automatization

Customization

Secure executions

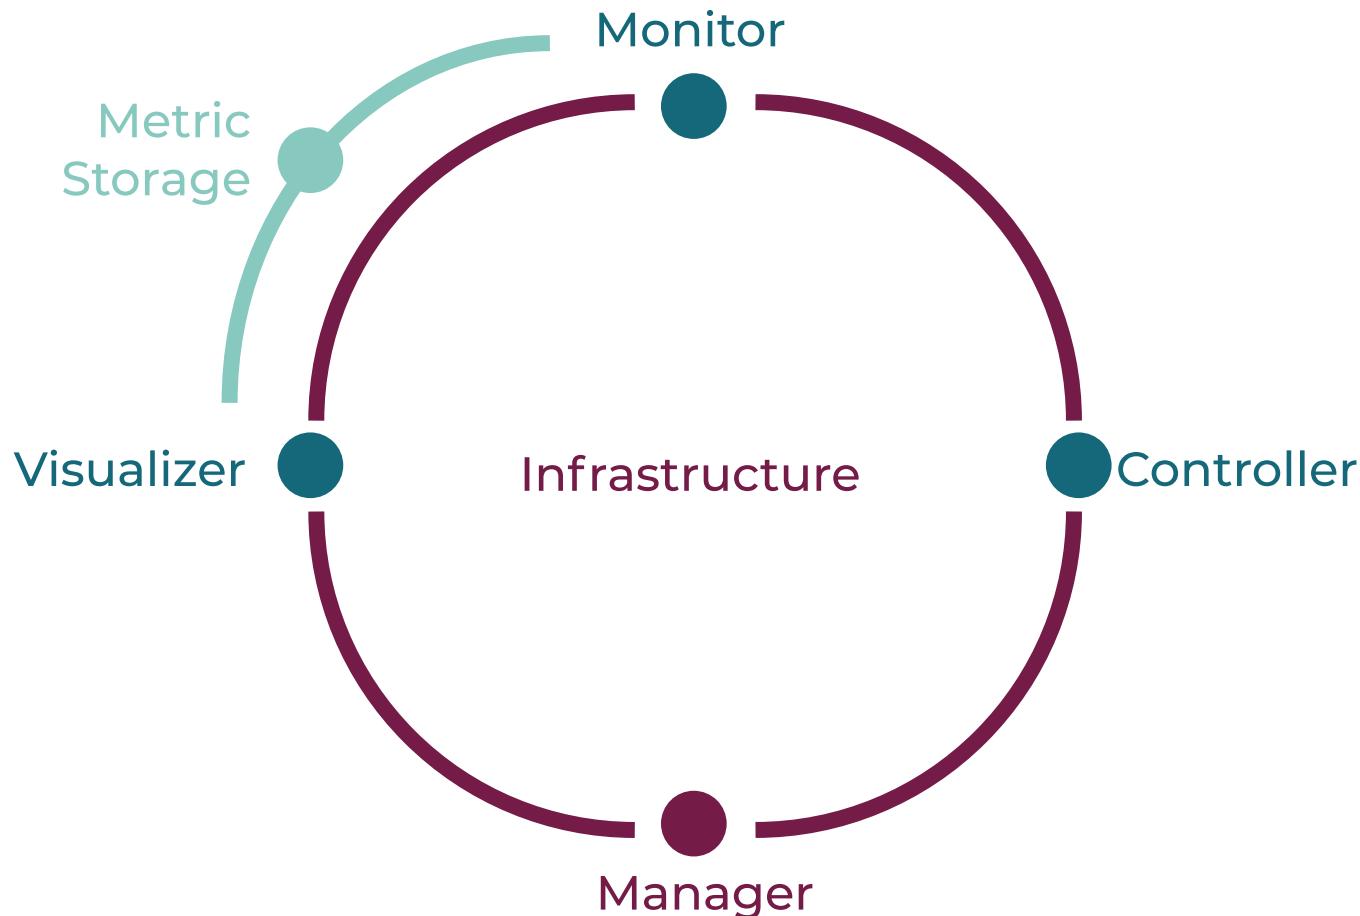


Figure 1. Asperathos architecture

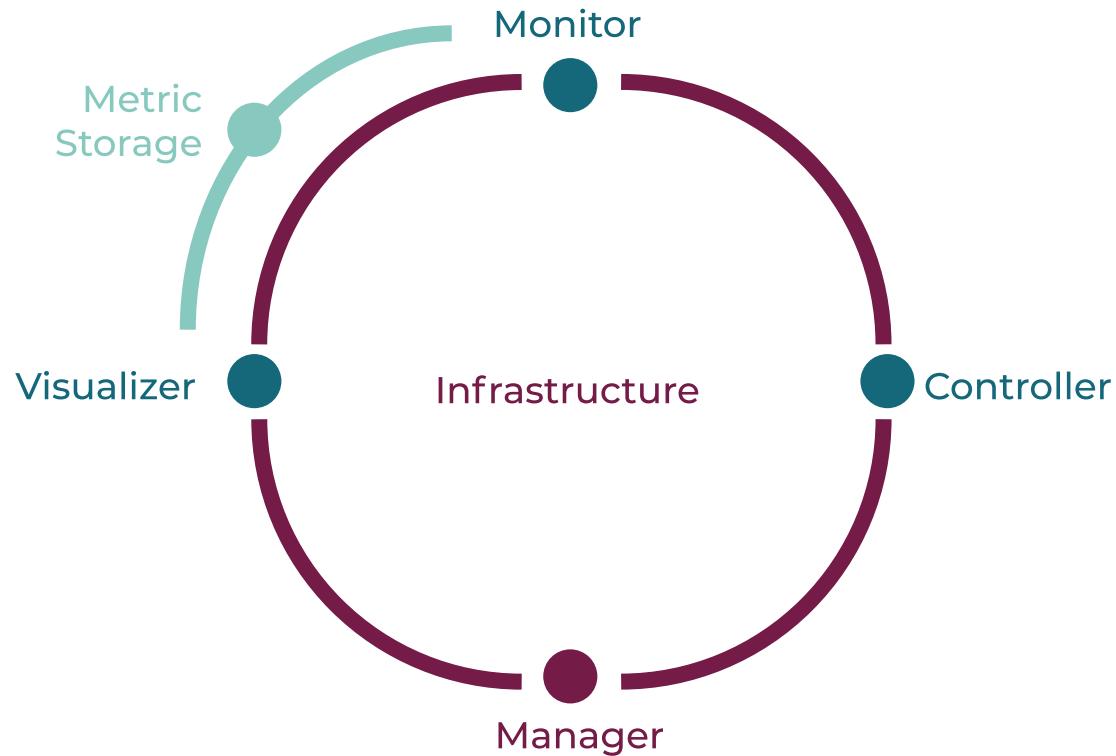


Figure 1. Asperathos architecture

Controlling the system in order to meet deadlines can be difficult

What can Asperathos do?



# Confidential data processing



# QoS-aware data processing



# Confidential data processing



# QoS-aware data processing

# Using SCONE to build SGX applications

# Intel SGX In Its Original Design

---

**Intention:** Only for very small functionality like generating secrets

**Complicated usage:** sgx\_create\_enclave

System call interface access through e-calls & o-calls

```
enclave_ref(enclave);
while (!enclave_terminated(enclave)) {
    enclave_enter(enclave, args->tcs_id, args->call_id, args, ret);
    switch (ret[0]) {
        case EXIT_TERMINATE:
            exit_code = ret[1];
            enclave_terminate(enclave);
            exit(exit_code);
        case EXIT_CPUID: {
            unsigned int* reg = (unsigned int*)ret[1];
            do_cpuid(reg);
            args->call_id = ECALL_SYSCALL_RESUME;
            break;
        }
        case EXIT_SYSCALL: {
            syscall_t* sc = (syscall_t*)ret[1];
            if (sc->syscallno == SYS_vfork){
                pid_t pid = vfork();
                sc->syscallno = pid;
            } else {
                scall(sc);
            }
            args->call_id = ECALL_SYSCALL_RESUME;
            break;
        }
        case EXIT_SLEEP: {
            __atomic_store_n(&E->ethread_state[args->id], ETHREAD_SLEEP)
```

# SCONE'S Design Goals

---

Minimal developer effort: Compile w/ **scone-gcc** instead w/**gcc**

- Alternatively, use prebuilt scone docker images

Run entire application in enclave

Provide transparent attestation, encryption and secret injection (Palaemon)

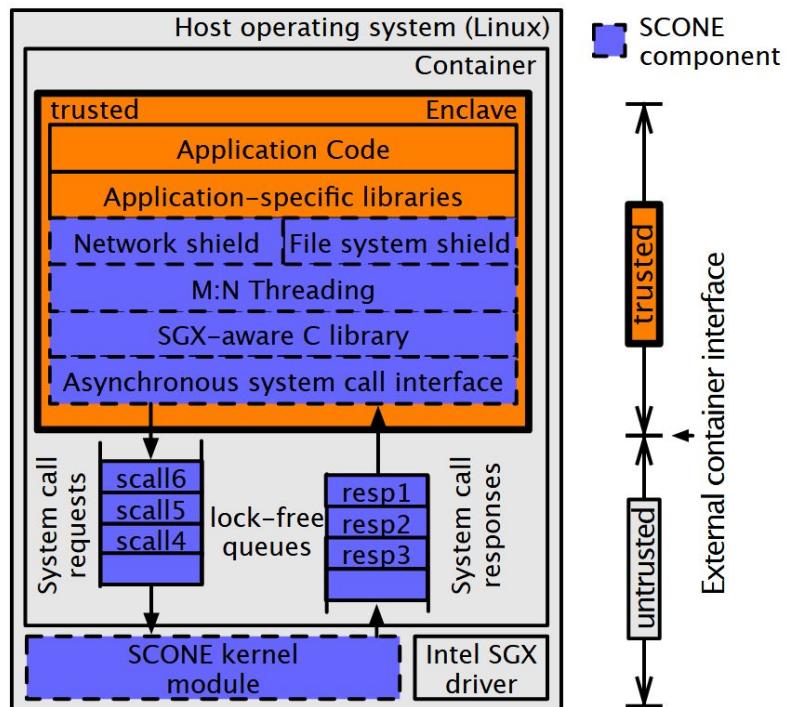
Tight integration in eco-systems, i.e., Docker & Swarm, Kubernetes

# SCONE Under The Hood

Starter code

System call interface

User level scheduling



# What is SCONE?

---

- 1) **Cross Compiler** to “sconify” applications, i.e., run them in Intel SGX enclaves
- 2) A **System Library** to provide system call support to talk to the external world

# How To Use SCONE? 5 Easy Steps

---

- 1) **Enable SGX in Bios** (if not done already)
- 2) **Install Intel SGX Drivers**
- 3) **Download/pull cross compiler docker image**
- 4) **Compile** your favorite application
- 5) **Run** your application

# How To Use SCONE? Step #1 - Enable Intel SGX in Bios

---

Under **Security** -> **Intel SGX**

Usually three options:

1. **Disabled**
2. **Enabled <- to choose**
3. **Software controlled**

# How To Use SCONE? Step #2 - Install Intel SGX Drivers

---

Use the following one liner:

```
$ curl -fssl https://tinyurl.com/y2byyh4h | bash
```

Or follow official steps:

<https://github.com/intel/linux-sgx-driver#install-the-intel-sgx-driver>

# How To Use SCONE? Step #3 - Download cross compiler docker image

---

Use the following two one liners:

```
$ docker pull sconecuratedimages/issre2019:crosscompilers
```

(This is the SCONE cross-compiler image for scone-based compilation based on the Alpine Linux docker imager)

```
$ docker pull alpine
```

(This is the bare bone Alpine Linux docker image for native compilation)

# How To Use SCONE? Step #5 - Compile your favorite application

---

```
#include <iostream>
#include <cmath>

using namespace std;

int main() {
    char* secret = (char*)"Karate";
    int x = 0;
    while(x < 10) {
        double y = sqrt((double)x);
        cout << "The square root of " << x << " is " << y << endl;
        x++;
    }
    cout << secret << endl;
    do cout << '\n' << "Press a key to continue...";
    while (cin.get() != '\n');
    return 0;
}
```

# How To Use SCONE? Step #5 - Compile your favorite application

---

```
$ wget -O sqrt.cc https://tinyurl.com/y6nyt4ly
```

```
$ docker run -v $(pwd):/myApp --device=/dev/isgx -it  
sconecuratedimages/issre2019:crosscompilers
```

```
$ cd /myApp
```

```
$ g++ -o sqrt-scone sqrt.cc
```

# How To Use SCONE? Step #5 - Run your favorite application

— — —

```
$ SCONE_VERSION=1 ./sqrt-scone
```

That's it!

# Now We Do A Memory Dump (in a second terminal)

— — —

```
$ wget -O dump-memory.py https://tinyurl.com/y2x4nnyx
$ wget -O memory-dump.sh https://tinyurl.com/y3c6ucmw
$ chmod +x *.sh *.py
$ sudo ./memory-dump.sh

$ cat content-memory | grep Karate
```

# Now The Same Without SCONE And Compare

---

```
$ docker run -v $(pwd):/myApp -it alpine
```

```
$ cd /myApp && apk add g++
```

```
$ g++ -o sqrt-native sqrt.cc
```

```
$ ./sqrt-native
```

# Use case analysis: anonymization of sensitive echocardio

# The Radiomics application

Anonymizing sensitive echocadio data

- Sensitive information is removed from video frames
- 2 types of input
  - Default video by video
  - Video archives



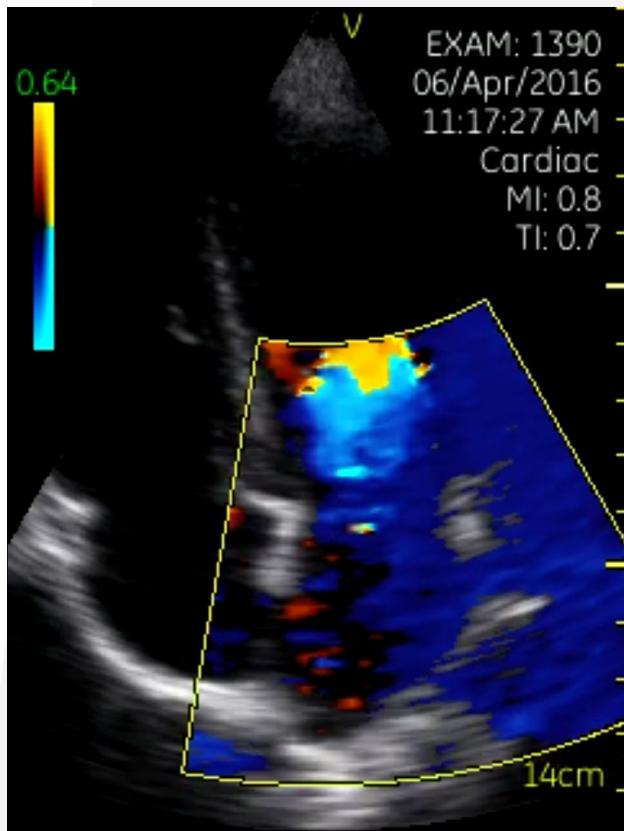


Figure 2. Radiomics video entry

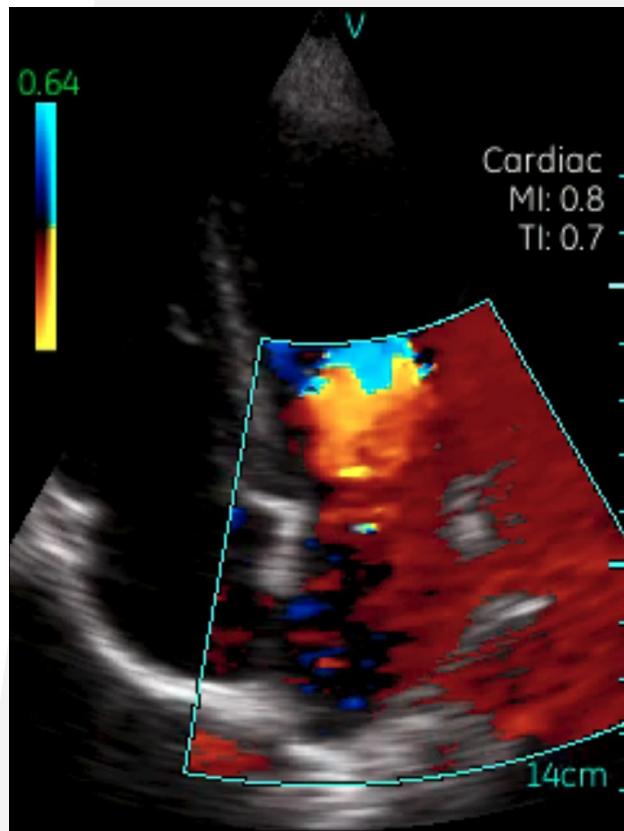


Figure 3. Radiomics anonymized result

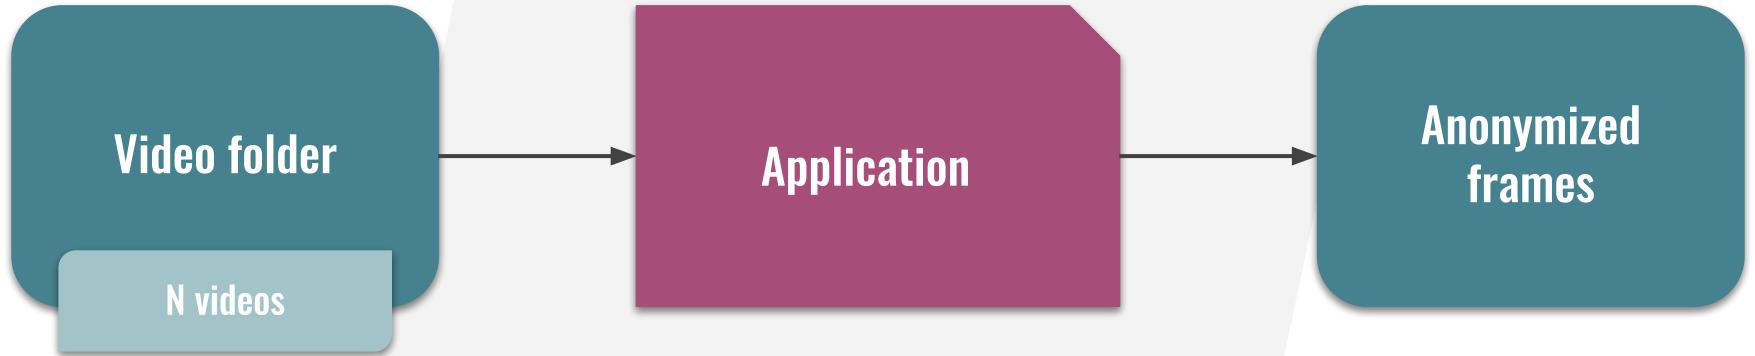


Figure 4. Radiomics simple architecture

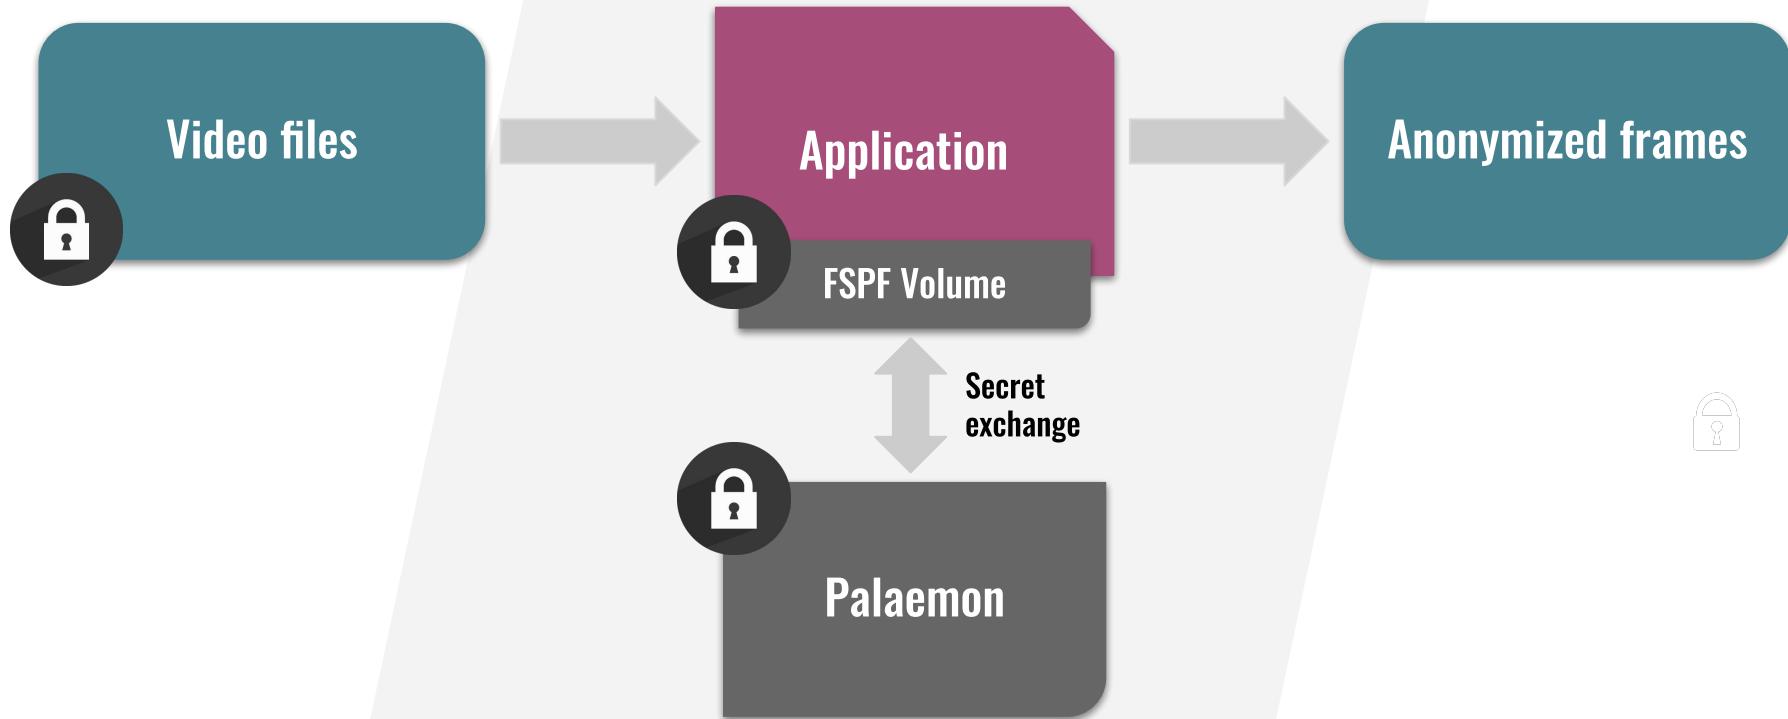


Figure 5. Radiomics architecture using SCONE and FSPF

# Performance Overheads - 1

Understanding the performance of the use case: Execution time for Radiomics using SCONE and FSPF

- **Scenarios**
  - Unprotected
  - Protected execution
  - Protected execution and FSPF
- **Factors**
  - Sample size
- **EPC size: 90MB**
- **Machine used**
  - Intel(R) Core(TM) i7-6700 CPU @ 3.40GHz
  - 16GB RAM



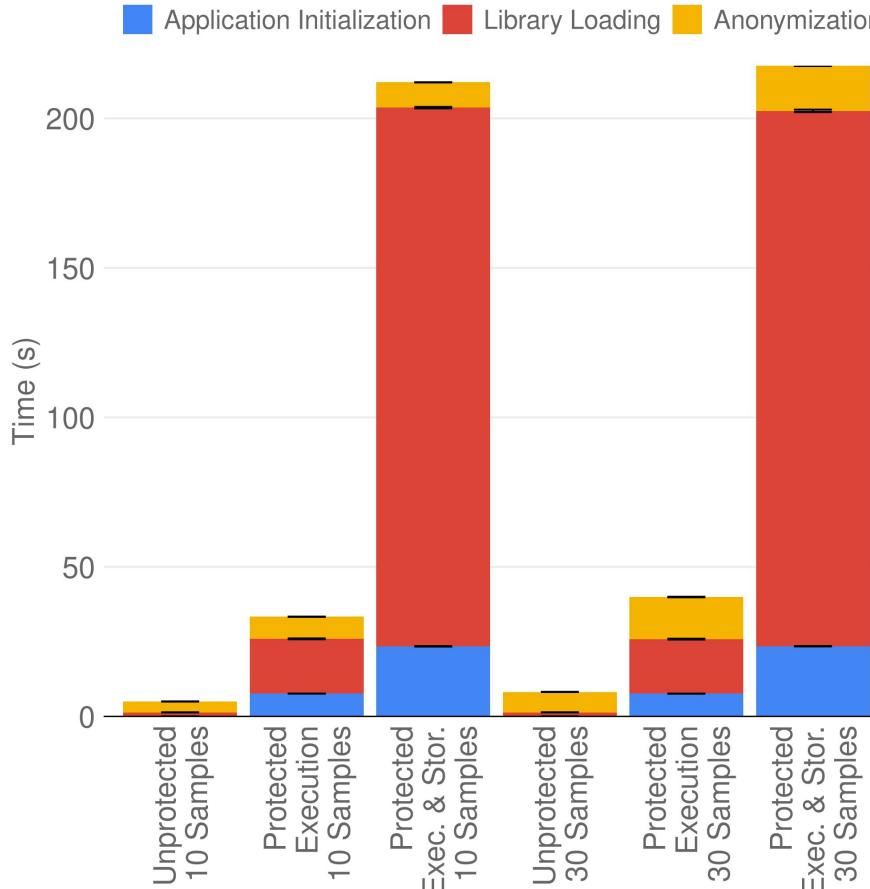


Figure 7. Experiment results considering execution time for SCONE executions

# Performance Overheads - 2

Understanding the performance of the use case: Execution time for Radiomics using SCONE and FSPF

- **Scenarios**
  - Unprotected
  - Protected execution
  - Protected execution and FSPF
- **Factors**
  - Sample size
  - **EPC size**
  - **Number of vCPUs**
- **Machine used**
  - Intel(R) Core(TM) i7-6700 CPU @ 3.40GHz
  - 16GB RAM



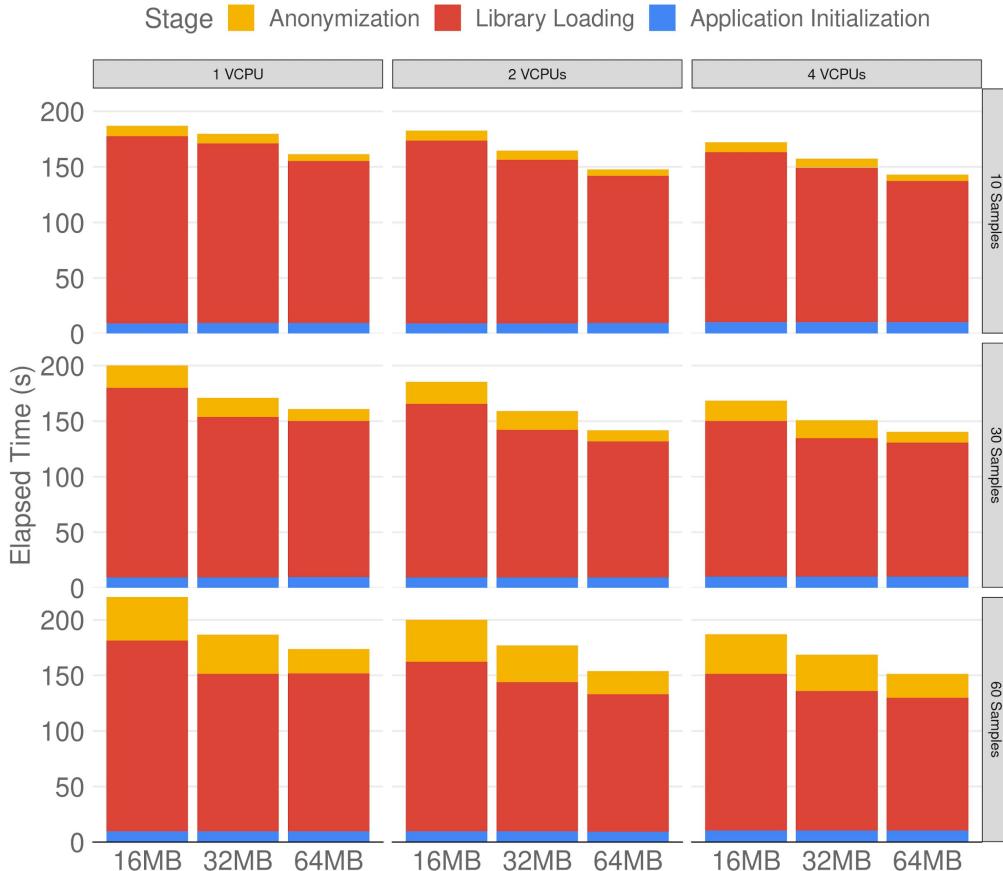


Figure 8. Experiment results considering execution time for SCONE executions varying EPC

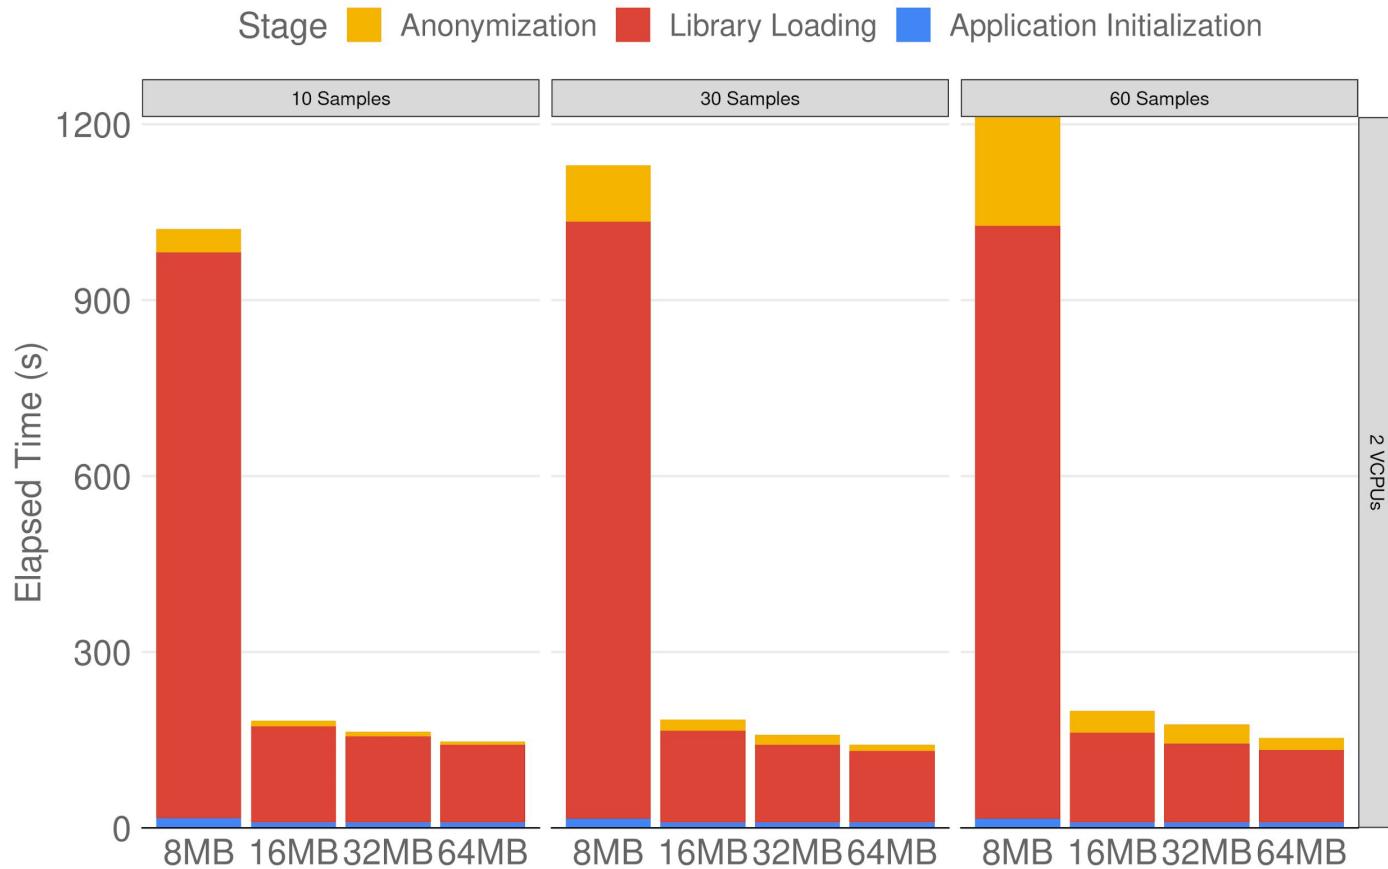


Figure 9. Experiment results show that for 8MB of EPC, elapsed time is much higher

# Lessons learned

What can we learn from these experiments?

- Requiring many third party libraries is expensive because they need to be integrity checked
- For batch processing multiple items at once may be desirable
- EPC can make a huge difference in some cases, (our 8 MB EPC example)
- It can be better to use four 1-CPU, 16MB EPC machines than one four-CPU, 64MB EPC machine



# Action! What do you need to perform an execution?

---

- The repositories used in this tutorial are available on GitHub
  - <https://github.com/ufcg-lsd/issre-tutorial>
- We are now going to perform 2 example executions:
  - SCONE + Radiomics
  - SCONE + FSPF + Radiomics
  - Reference to the guide:
    - <https://github.com/ufcg-lsd/radiomics-scone>

# QoS and security-aware data processing

Part 2



# Confidential data processing



# QoS-aware data processing

# Kubernetes 101

# Nodes, clusters and volumes

The hardware

- **Nodes**
  - Representation of a single machine in a cluster
  - Physical machines or virtual machines
- **Clusters**
  - Composition of a set of nodes
  - It shouldn't matter to the program which individual machines are actually running the code
- **Volumes**
  - Data can't be saved to any arbitrary place in the file system
  - Can be mounted to the cluster and accessed by containers in a given pod



# Containers, pods and deployments

The software

- **Containers**
  - Programs running in Kubernetes
  - Self-contained Linux execution environments
- **Pods**
  - Composition of a set of containers
  - Pods are used as the unit of replication in Kubernetes
- **Deployments**
  - Manages and monitor a set of pods
  - Declares how many replicas of a pod should be running at a time

— — —

# Why Kubernetes?

---

- Allow for the **deployment of jobs**!
- Jobs are good for batch processing
  - Processes that run for a certain time to completion
- Applications are monitored and taken care of
- Large and active community
- Supported by all major cloud providers

# Asperathos 101

## MANAGER

- ❑ Entry point to the users
- ❑ Receives a job submission and prepares its execution
- ❑ Handles cluster configuration
- ❑ Sends requests to the Monitor and Controller components

## CONTROLLER

- ❑ Adjusts the amount of resources allocated to an application in order to guarantee QoS

## Monitor

### Metric Storage

### Visualizer

## Infrastructure

### Controller

### Manager

## MONITOR

- ❑ Calculates and publishes application metrics (M1) or resource metrics (M2)
- ❑ M1. Application progress; M2. CPU, memory, etc.

## VISUALIZER

- ❑ Allows the visualization of a job's progress
- ❑ Consumes metrics from the Monitor and generates graphs
- ❑ Grafana
- ❑ Influxdb and Monasca

# What is a job?

- **"plugin": "kubeflow"**
- **"cmd": ["python", "app.py"]**
- **"img": "image/name"**
- **"init\_size": 1**
- **"redis\_workload":**  
["https://gist.github.com/username/raw/43eaeffe10"](https://gist.github.com/username/raw/43eaeffe10)
- **"job\_resources\_lifetime": 30**
- **"control\_plugin": "kubeflow"**
- **"control\_parameters":**
  - **"schedule\_strategy": "default"**
  - **"actuator": "k8s\_replicas"**
  - **"check\_interval": 10**
  - **"trigger\_down": 1**
  - **"trigger\_up": 1**
  - **"min\_rep": 1**
  - **"max\_rep": 10**
  - **"actuation\_size": 1**
  - **"metric\_source": "redis"**
- **"monitor\_plugin": "kubeflow"**
- **"expected\_time": 130**
- **"enable\_visualizer": true**
- **"visualizer\_plugin": "k8s-grafana"**
- **"env\_vars": {}**

# What you need to use Asperathos?

Have a Kubernetes cluster ready to receive applications



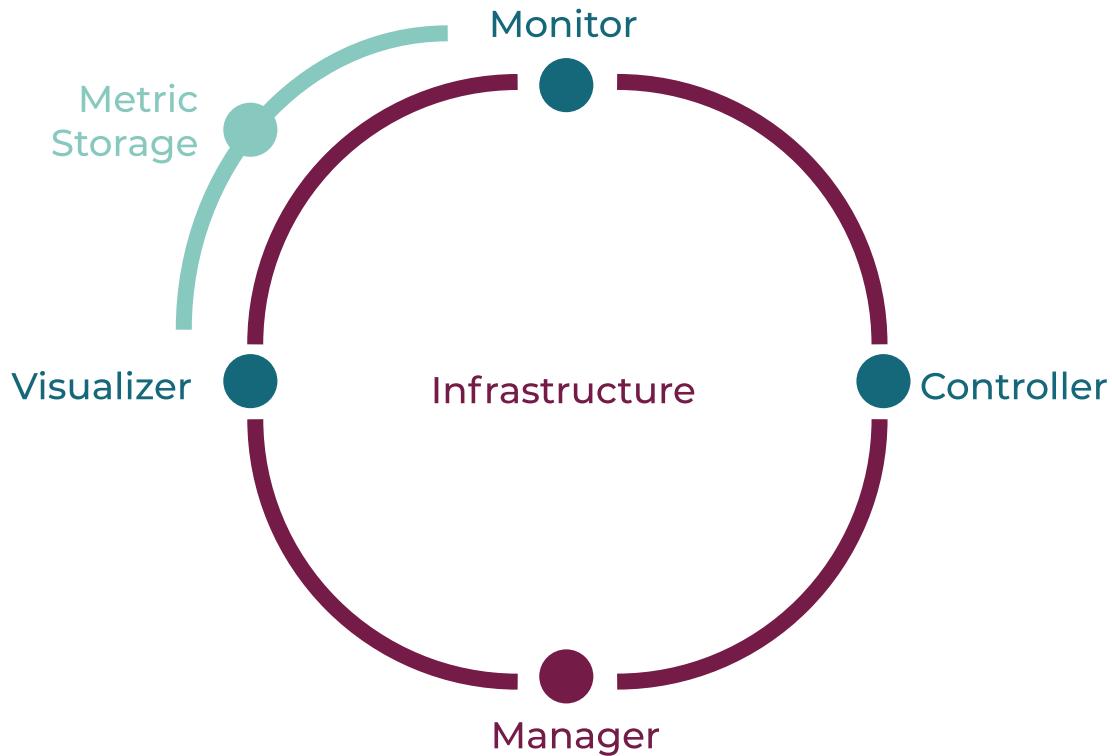
Have an application image capable consuming items from a redis workload

# Action! How to deploy an Asperathos instance?

---

- The repository for Asperathos and its components is available on GitHub
  - <https://github.com/ufcg-lsd/asperathos>
- We are now going to **deploy an instance of Asperathos**
  - Reference to the QuickStart guide in the GitHub page linked above

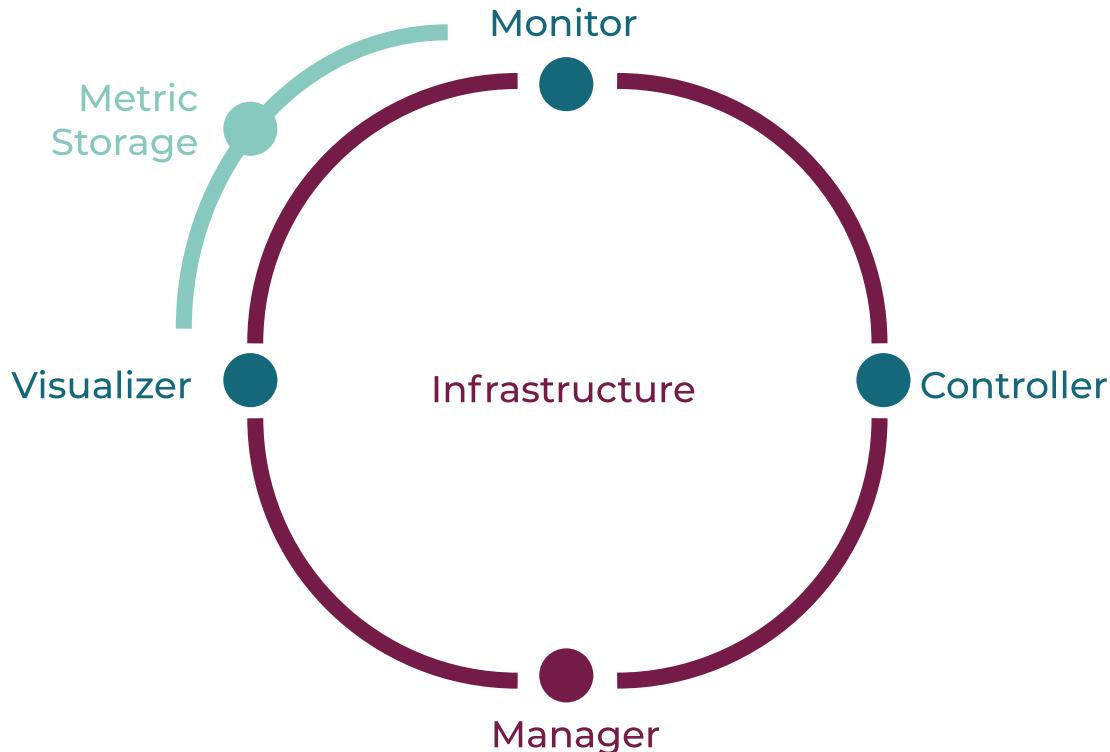
# Customizing Asperathos components



Asperathos  
architecture  
allows the addition  
of customized  
plugins

What can you customize?

# Controlling the resources used by a job



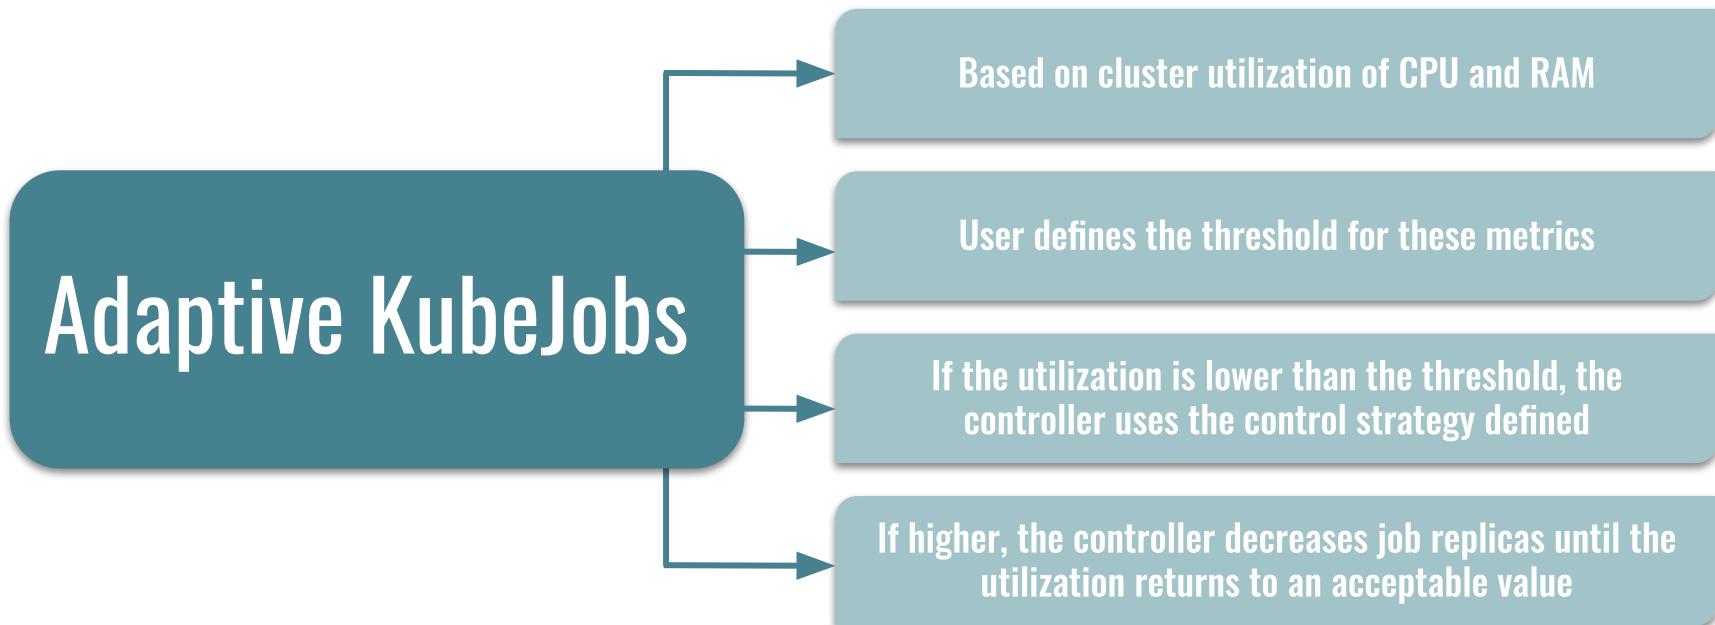
## Controller

- ❑ Responsible for adjusting the quantity of resources allocated to the execution of one application in a way that deadlines are met and QoS guaranteed

## Customizable strategies

- ❑ Default: fixed actuation size
- ❑ PID: actuation size depends on the error magnitude and trend
- ❑ Your own strategy!

# Creating a plugin for the Controller component



# How to install

Send a POST request to <your\_asperathos\_url:port>/plugins with this json body:

```
{  
    "plugin_source":  
    "https://git.lsd.ufcg.edu.br/asperathos-custom/adaptive-kubejobs",  
    "install_source": "git",  
    "plugin_module": "adaptive_kubejobs",  
    "component": "controller",  
    "plugin_name": "adaptive_kubejobs"  
}
```

# How to use

**Requirements:** metric-server by Kubernetes

Json must contain the parameters bellow:

```
"control_plugin": "adaptive_kubejobs",
"control_parameters": {
    "schedule_strategy": "default",
    "actuator": "k8s_replicas",
    "check_interval": 10,
    "trigger_down": 1,
    "trigger_up": 1,
    "min_rep": 1,
    "max_rep": 10,
    "actuation_size": 1,
    "metric_source": "redis",
    "max_rammax_cpu
```



**Confidential** + **QoS-aware data**  
**data processing** **processing**

# Use case analysis: combining security and QoS

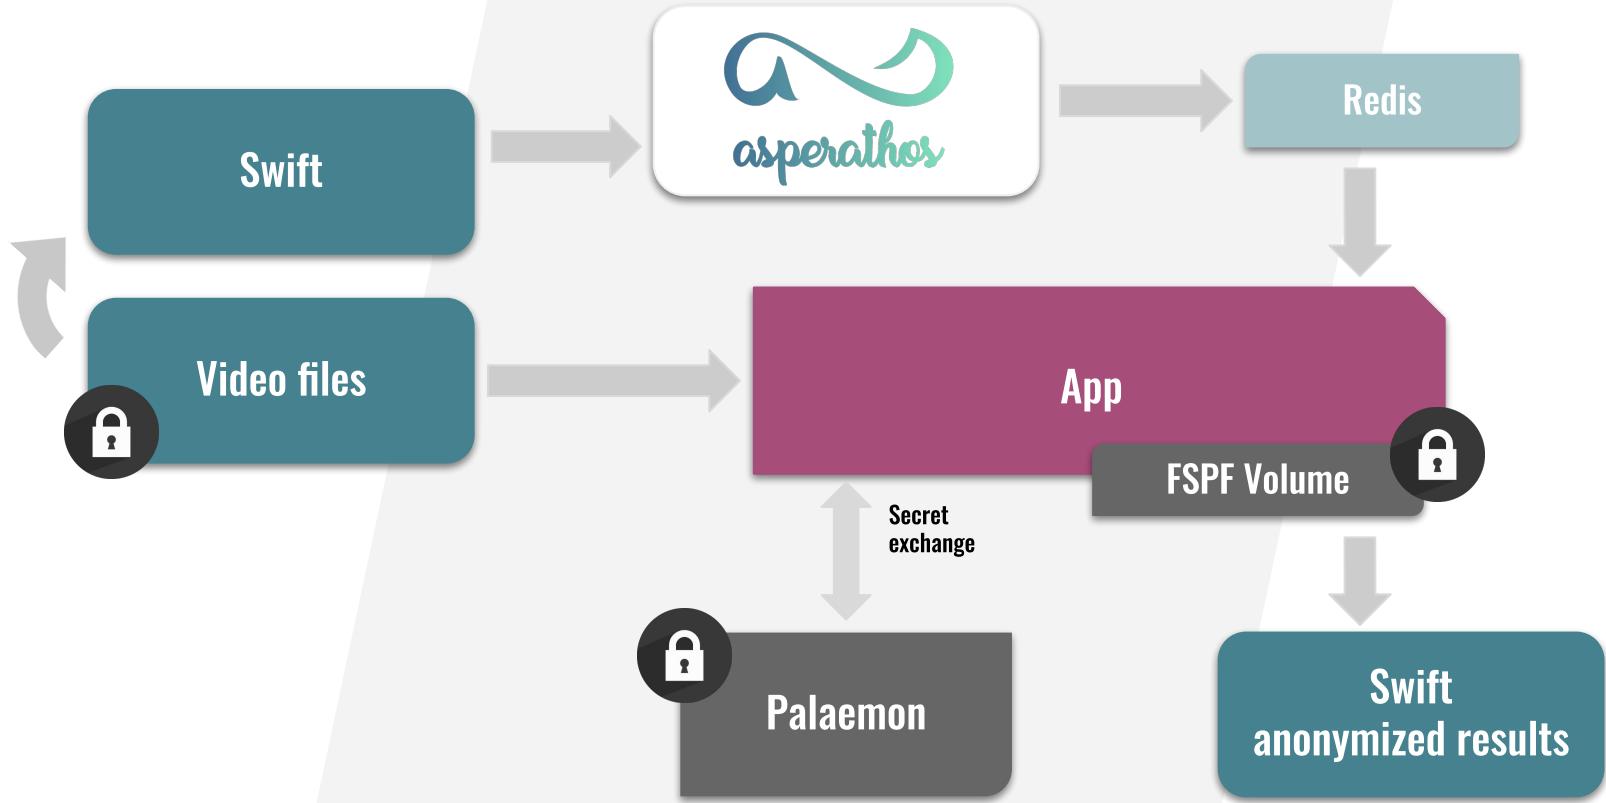


Figure 9. Architecture of the combination of Asperathos and Radiomics

# Action! What do you need to perform an execution?

---

- The repositories used in this tutorial are available on GitHub
  - <https://github.com/ufcg-lsd/issre-tutorial>
- We are now going to perform an execution of Radiomics using Asperathos
  - Reference to the guide:
    - <https://github.com/ufcg-lsd/radiomics-asperathos>

# Conclusions

# Conclusions

---

- SGX is a promising technology that enables sensitive data to be processed with confidentiality and integrity guarantees in untrusted clouds
- SCONE provides transparency to some non-trivial aspects of SGX programming, such as remote attestation and file/network encryption
- If application's working memory fits in the EPC available in your VM/machine, processing overhead is small
- But it is also important to minimize the number of files that need to be authenticated for the application
- If lots of data needs to be confidentially processed, Asperathos is a good alternative for orchestrating batch executions with QoS

Thank you for attending!

## **ISSRE 2019**

### Tutorial Session

**If you want to provide  
some feedback or have  
questions, please  
contact us!**

# **Building Applications for Trustworthy Data Analysis in the Cloud**

**Andrey Brito**

[andrey@computacao.ufcg.edu.br](mailto:andrey@computacao.ufcg.edu.br)

**André Martin**

[andre.martin@tu-dresden.de](mailto:andre.martin@tu-dresden.de)

**Lilia Sampaio**

[liliars@lsd.ufcg.edu.br](mailto:liliars@lsd.ufcg.edu.br)

**Fábio Silva**

[fabiosilva@lsd.ufcg.edu.br](mailto:fabiosilva@lsd.ufcg.edu.br)